

# Research report : renaming in Identifier-based Sequence Conflict-free Replicated Data Types (CRDTs)

Matthieu Nicolas

December 8, 2017

## 1 Context

### 1.1 System model

- Distributed large-scale system
- Asynchronous network
- Partition-tolerant
- Replicated sequence among nodes
- Eventual consistency
- Use a Identifier-based Sequence CRDT as the conflict resolution mechanism
- Intention preserving

### 1.2 Identifier-based Sequence Conflict-free Replicated Data Types (CRDTs)

#### 1.2.1 State

Has a state  $S$  which represents the replicated sequence (use additional metadata to do so)

- Noted as  $[(id, elt)]$  in the following figures
- The function  $view(S)$  allows to retrieve the sequence represented by the state  $S$
- **Example:**  $view([(id_1, elt_1), (id_2, elt_2)]) = [elt_1, elt_2]$

#### 1.2.2 Identifiers

Associates an identifier  $id$  to each element  $elt$  of the sequence

- Unique (an identifier can not be generated twice)
- Order relation (so that we can compare two identifiers)
  - Allows to determine the order of elements of the sequence using their identifiers

- Belong to a dense set
  - Always able to add a new element (and thus a new identifier) between two other elements

The elements in the sequence are always ordered according to their identifiers : in a sequence  $[(id_1, elt_1), \dots, (id_3, elt_3), \dots, (id_2, elt_2)]$  we always have  $id_1 < \dots < id_3 < \dots < id_2$ .

### 1.2.3 Operations

For each operation to update the data structure, has two forms of it: the *local* form and the *remote* one

- The *local* operation is triggered by the node (by user request for example)
- Performing a *local* operation on a given state  $S$  returns the new state  $S'$  and the metadata needed to build an equivalent *remote* operation
- The *remote* operation is propagated to other nodes so they can also update their own state
- Given a state  $S$  and an operation  $local(S, data) = (S', metadata)$ , we have  $remote(S, metadata) = S'$
- **Note:** given an *local* operation  $localOp$ , there may be several equivalent *remote* operations  $remoteOp, remoteOp', remoteOp'' \dots$

### 1.2.4 add

The operation *add* allows to insert an element into the sequence :

- $addLocal(S, index, elt) = (S', (id, elt))$ 
  - Update state  $S$  by adding an element  $elt$  at the position  $index$  in the sequence
  - Return the resulting state  $S'$  as well as the identifier  $id$  generated for this element
  - The identifier  $id$  will be generated according to the identifiers of the elements previously at the positions  $index - 1$  and  $index$ 
    - \* **Example:**  $addLocal([(id_1, elt_1), (id_2, elt_2)], 1, elt_3)$  will return  $id_3$  such as  $id_1 < id_3 < id_2$
  - This identifier  $id$  will be used (and especially its order relation with other identifiers) to update correctly other nodes' state
  - **Note:** When generating a new identifier between  $id_1$  and  $id_2$ , there may be several identifiers  $id_3, id'_3, id_3'' \dots$  such as  $id_1 < id_3 < id'_3 < id_3'' < id_2$ . The returned identifier is chosen in a undeterministic manner.
- $addRemote(S, id, elt) = (S', (index, elt))$ 
  - Update state  $S$  by adding an element  $elt$  in the sequence
  - The position of insertion of this element will be determined using its  $id$
  - Return the resulting state  $S'$  as well as the current index of the element in the sequence
- Given a state  $S$ , to one *addLocal* operation on  $S$ , many *addRemote* correspond (since the resulting  $id$  is generated in an undeterministic manner)
- Given a state  $S$ , to one *addRemote* operation on  $S$ , only one *addLocal* corresponds

### 1.2.5 *del*

The operation *del* allows to remove an element from the sequence :

- $delLocal(S, index) = (S', id)$ 
  - Update state  $S$  by removing the element at the position  $index$  in the sequence
  - Return the resulting state  $S'$  as well as the identifier  $id$  of the deleted element
- $delRemote(S, id) = (S', index)$  allowing to remove the element identified by  $id$ 
  - Update state  $S$  by removing the element identified by  $id$
  - Return the resulting state  $S'$  as well as the position  $index$  of the deleted element in the sequence
- Given a state  $S$ , to one *delLocal* operation, only one *delRemote* corresponds
- Given a state  $S$ , to one *delRemote* operation, only one *delLocal* corresponds

### 1.2.6 Log of operations

Associates to a state  $S$  a log  $L$

- Is a sequence of entries  $(remoteOp, localOp)$ , a remote operation and its local counterpart
- The sequence of remote operations, performed in order from a blank state  $S_{blank}$ , allows to recreate state  $S$
- Each entry represented as 

$remoteOp$
$localOp$

 in the following figures

### 1.2.7 Causal context

Associates to a state  $S$  a causal context  $cc$

- Represents all operations known at state  $S$
- Can use a *version vector* for example as an implementation

An example of the lifecycle of such a replicated data structure is shown in figure 1

## 2 *rename* operations

### 2.1 Motivation

- Identifiers growing over time
- Performances of the data structure thus decreasing over time

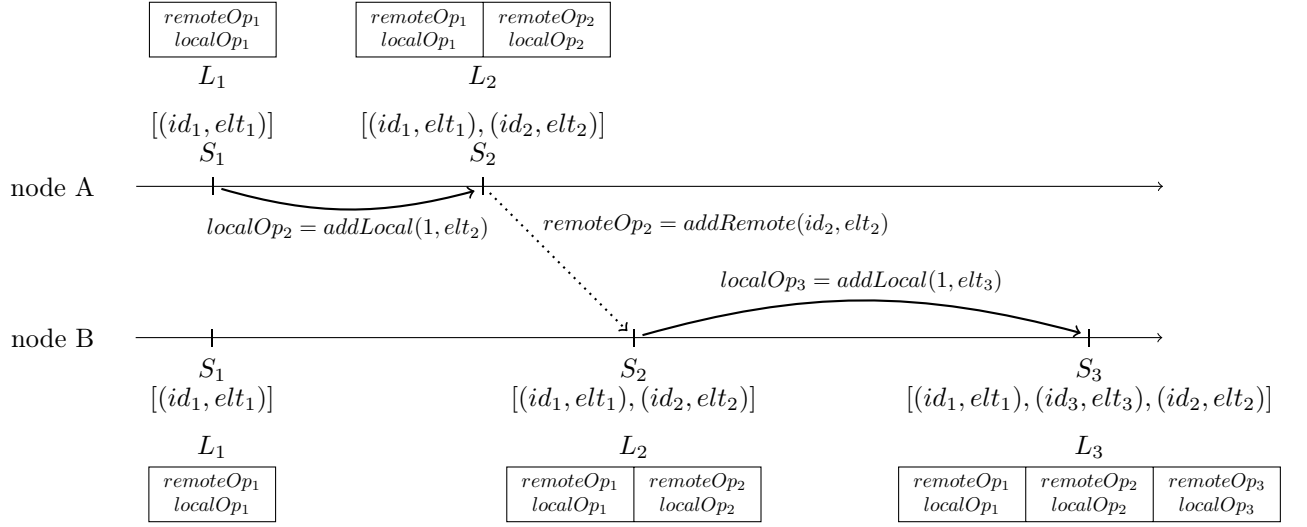


Figure 1: Insertion of elements in the replicated sequence

## 2.2 *renameLocal*

- Add an operation  $renameLocal(S) = (S', mapIds, cc_S)$ 
  - Replace each identifier attached to elements of  $S$  with new ones
  - Return a map  $mapIds$  of the previous identifiers to the new ones
  - Also need to return the causal context  $cc_S$  of the state  $S$  to indicate on which state has been performed the renaming operation
  - $view(S) = view(S')$  where  $(S', \_, \_) = renameLocal(S)$
  - Represented by figure 2




Figure 2: Local renaming of identifiers of the replicated sequence

## 2.3 *renameRemote*

- Add an operation  $renameRemote(S, L, mapIds, cc_{S'}) = (S'', L'')$

- Replace current state  $S$  by equivalent state  $S''$  and current log  $L$  by equivalent log  $L''$
- Rename all identifiers  $id \in S \cdot id \in S'$  using  $mapIds$
- Also have to rename all identifiers  $id \in S \cdot id \notin S'$  to preserve the current order of elements
- **Precondition:**  $S \geq S'$  ( $S$  has seen all the operations seen by  $S'$  but may have seen more)
- $view(S) = view(S'')$  where  $(S'', \_) = renameRemote(S, L, mapIds, cc_{S'})$

## 2.4 Usage

Given an operation  $renameRemote(S, L, mapIds, cc_{S'})$ , resulting from the execution of  $renameLocal(S')$  on another node, we have to perform the following steps to apply it:

1. Instantiate a blank state  $S''$  and its empty log  $L''$
2. Generate a log  $L_{causal}$  made of all operations belonging to the causal context  $cc_{S'}$
3. For each entry  $(remoteOp, localOp)$  of  $L_{causal}$ 
  - (a) Update state  $S''$  by performing  $remoteOp(S'', metadata)$
  - (b) Add entry  $(remoteOp, localOp)$  to  $L''$
4. Rename all identifiers of the data structure according to  $mapIds$  (at this point,  $S'' = S'$ )
5. Generate a log  $L_{concurrent}$  made of all operations of  $L$  not included in  $L_{causal}$
6. For each entry  $(remoteOp, localOp)$  of  $L_{concurrent}$ 
  - (a) Update state  $S''$  by performing  $localOp(S''_{prev}) = (S''_{new}, metadata)$
  - (b) Build new remote operation  $remoteOp'$  given  $metadata$
  - (c) Add entry  $(remoteOp', localOp)$  to  $L''$
  - (d) Propagate  $remoteOp'$

This algorithm is represented by figure 3

## 2.5 Limits

- Different nodes, while performing the remote renaming operation, may replay at step 6a the same operation
- Since there is no coordination between them, in the case of a  $addLocal$ , they will end up generating two different remote operations  $remoteOp'$  and  $remoteOp''$  during step 6b
- We will have to deliver them both to each node to actually converge (the states would differ otherwise)
- This will result in the duplication of the user's intention (since the inserted element will end up being added twice)
- An example is shown in figure 4




Figure 3: Renaming with concurrent operations

### 3 *addRedo* operation

#### 3.1 Idea

- At step 6a, if we can generate deterministically the resulting *id* for a given previous log entry (*addRemote*, *addLocal*), then we would not duplicate the user's intention
- Indeed, each node would thus generates the same operation *addRemote'* at step 6b
- In that case, we would only need to deliver at least once *addRemote'* to the nodes to converge (or exactly-once if the *addRemote* is not idempotent)




Figure 4: Duplication of the intention of  $localOp_3$

### 3.2 Research question

Can we define the following operation  $addRedo(S, (addRemote, addLocal)) = (S', (id', elt))$  such as :

- $id'$  is generated deterministically
- $view(S) = view(S')$  where  $(S, id) = addLocal(S'', index, elt)$  and  $(S', \_) = addRedo(S'', (addRemote(S'', id, elt), addLocal(S'', index, elt)))$

This operation would be used at step 6a instead of simply using  $addLocal$  and would solve the duplication effect.

## 4 Discussion

- Need to keep the log of operations (both *remote* and *local*)
- Performances of a *renameRemote* depend on the number of operations in the log and the number of concurrent operations
  - Have to replay all operations from the causal context of the *renameRemote* operation
  - Have to regenerate concurrent operations and propagate them
- Can propose mechanism to reduce the size of the log
  - By pruning causally stable entries and using snapshots

- New identifiers generated by *addRedo* operations may be larger than the initial ones according to the chosen strategy
  - Can argue that they will shrink with the next *rename*
- Solving concurrent *rename* looks difficult
  - For now, can assume that only one node can perform such operations

## 5 Counter-example

Found a counter-example which invalidate the algorithm proposed in section 2.4. We replay the same operations as in previous examples but in this case:

- *localOp<sub>2</sub>* and *localOp<sub>3</sub>* are concurrent
- The generated identifiers *id<sub>2</sub>* and *id<sub>3</sub>* are in this order:  $id_2 < id_3$

In this scenario, when replaying *localOp<sub>3</sub>* at step 6a, we will swap the position of the elements *elt<sub>2</sub>* and *elt<sub>3</sub>* compared to the previously observed state. This result in a incoherence of the system and may result in the violation of the intention of following operations based on this previously observed state. This scenario is represented by figure 5.



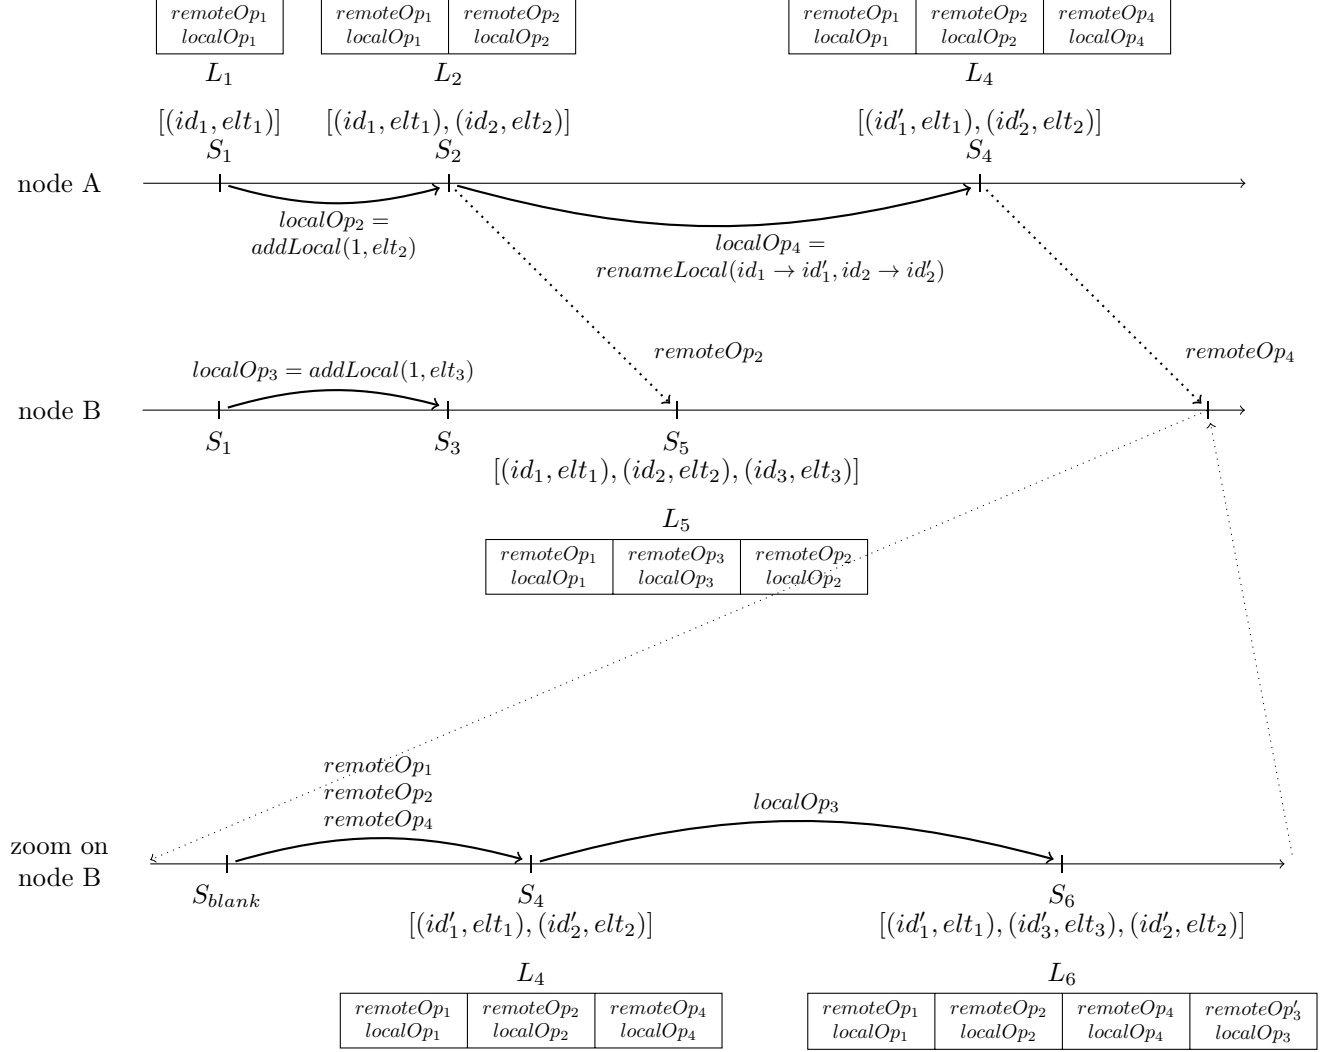


Figure 5: Incoherence occurring by replaying local operations during renaming process